

A598a: Machine Learning in Astronomy

Andrew Connolly and Stephen Portillo
University of Washington, Winter Quarter 2022

Class and Office hours

Class will be held Mondays and Wednesdays in PAB B305 from 2.00pm to 3.20 pm.

Office hours are on a drop-in basis and by appointment.

Class materials

Website and repository: <https://github.com/uw-astro/astr-598a-win22>

Class JupyterHub: <https://rttl.axdd.s.uw.edu/2022-winter-astr-598-a/>

Textbook: Ivezić, Connolly, VanderPlas & Gray “Statistics, Data Mining, and Machine Learning in Astronomy: A Practical Python Guide for the Analysis of Survey Data”

Goals and Objectives

This course will introduce graduate students to common statistical and machine learning methods used in astronomy and other physical sciences. While it will include theoretical and methodological backgrounds for the techniques, the focus will be on the *application* of machine learning to astrophysical problems. Practical data analyses will be done using python tools, such as astroML module (see www.astroML.org), scikit-learn, and tensor flow/pytorch and applied to astronomical datasets. While focused on astronomy, this course should be useful to all graduate students interested in data analysis in the physical sciences and engineering. The lectures will be aimed at graduate students and the main discussion topics will be based on selected topics from Chapters 6-10, in the reference textbook “Statistics, Data Mining, and Machine Learning in Astronomy: A Practical Python Guide for the Analysis of Survey Data”.

The goal of this course is to give you the tools necessary to understand and analyze rich datasets, such as those from the SDSS to the LSST.

Homework and Final Project:

There will typically be one homework each week. Homework will focus on practical applications using Python, designed to exercise what we have learnt in the week prior to the homework being assigned. All homework will involve writing Jupyter notebooks and submitting them via git.

In place of a final exam, you will propose a project (a piece of software, or an analysis) to build or improve using the techniques and libraries we'll learn about in the course. Ideally, this is something that helps with your research. This project can continue from the project you started in the "Introduction to Astrostatistics and Data-Intensive Astronomy" class.

Prerequisites:

Students taking this class are expected to have a background in calculus and Python and a basic understanding of statistics (e.g., Bayes theory and its application, and maximum likelihood estimation). Taking ASTR 598A: "Introduction to Astrostatistics and Data-Intensive Astronomy" (offered in the Autumn quarter) is strongly encouraged

Course Structure

Lecture 1	Introduction to the course and project overview
Lecture 2	Regression: Linear, Non-linear, Outliers, Errors on Variables
Lecture 3	Regression: Basis Function Regression, Regularization, Cross-Validation
Lecture 4	Dimensionality Reduction: PCA, Missing Data, NMF, ICA, LLE
Lecture 5	Time Series: Fourier Analysis, Filtering, Autoregressive processes, Correlation functions
Lecture 6	Density Estimation: K-means, Kernel Density Estimation, K-nearest neighbors
Lecture 7	Density Estimation: Gaussian Mixture Models, eXtreme Deconvolution
Lecture 8	Supervised Classification: Naïve Bayes, GMM, LDA, ROC Curves
Lecture 9	Supervised Classification: Decision Trees, Random Forests
Lecture 10	Hierarchical Bayes
Lecture 11	Approximate Bayesian Computation
Lecture 12	Deep Learning: Neural Networks
Lecture 13	Deep Learning: Convolutional Neural Networks
Lecture 14	Deep Learning: Autoencoders and Dimensionality Reduction
Lecture 15	Project time
Lecture 16	Project time
Lecture 17	Project Presentations

Note that these topics and timing may change dependent on the progression and interests of the class. Additional topics we may cover based on interest include: Variational Inference, Probabilistic Programming.